
Lost in Reasoning: Bistable CoT Faithfulness Across Multi-Turn Conversations and Its Implications for Safety Monitoring

Dhruv Gupta

BITS Pilani

f20221683@pilani.bits-pilani.ac.in

Abstract

When a reasoning model produces a chain-of-thought trace, we assume the trace reflects what the model is actually computing. That assumption underpins a growing class of AI safety proposals: if traces are faithful, auditors can read them to verify the model is reasoning correctly. But is the assumption warranted during long, multi-turn conversations where a model has already “made up its mind”?

We study this question by surgically removing a model’s reasoning at every turn of 67 multi-turn conversations and asking: does the model’s answer change? If erasing the reasoning changes nothing, the trace was decorative — the model had already committed to an answer and the chain-of-thought merely dressed it up.

We find that faithfulness is **not** a fixed property of a model: within a single conversation the same model alternates between two distinct modes. In *exploring* mode the reasoning genuinely drives the answer; in *anchored* mode the reasoning is post-hoc rationalization — erasing it makes no difference. Across 67 conversations and 1,289 turn-level measurements, 23.3% of turns are anchored.

Crucially, anchoring is not random noise: across our 67 conversations some conversations are heavily anchored while others barely anchor at all, a pattern confirmed statistically and replicated across three independent experimental seeds (anchoring 27–38% per seed). More worrying still, the fraction of anchored turns grows $3.9\times$ as conversations lengthen—from 7.2% in short conversations to 28.4% in long ones. This means CoT-based monitoring becomes least reliable precisely in the long, high-stakes settings where it is most needed.

1 Introduction

The faithfulness assumption in deployed AI. Reasoning models that produce visible chain-of-thought (CoT) traces are used in high-stakes conversational settings—medical question-answering, tutoring, coding assistance—on the premise that the trace reveals what the model is computing. This assumption matters for safety: if the trace is faithful, an auditor can inspect it to verify the model is reasoning from the right information; if not, the trace is uninformative or worse—a confident-looking rationalization of a predetermined answer. Existing faithfulness evaluations test only *single-turn* prompts [Lanham et al., 2023, Turpin et al., 2023, Arcuschin et al., 2025], leaving open whether faithfulness holds as a conversation unfolds across many turns.

The gap: multi-turn derailment and its effect on faithfulness. Laban et al. [2025] showed that reasoning models lose roughly 39 percentage points of accuracy when multi-turn problems are revealed piece by piece across turns, relative to single-turn baselines. The dominant failure mode is *premature commitment*: the model confidently states an early, often wrong, answer and does not

revise it as new information arrives. Laban et al. treat the CoT trace as a black box: they measure answer quality, not the relationship between the reasoning chain and the answer. The question left open is whether, during the “lost in conversation” phase, the model is still genuinely reasoning or merely rationalizing—producing a plausible-looking trace that does not constrain its output.

A novel finding: faithfulness is mode-switching, not monotonically decaying. Applying counterfactual deletion turn-by-turn to a 44-turn derailing conversation reveals that faithfulness does *not* decay monotonically with turn index. Instead, the model alternates between two quasi-stable modes. In **exploring** mode, truncating the CoT changes the answer—the reasoning chain is causally active. In **anchored** mode, all five truncation levels produce the same numeric answer regardless of how much reasoning is shown—the chain of thought is a post-hoc narrative for an answer the model had already committed to (Figure 1). The transition into anchored mode aligns with the onset of premature commitment, and the model can return to exploring mode when new information arrives.

Our method and results. We take 67 GSM8K sharded conversations (collected in seven independent datasets across two experimental phases) and, at every turn, surgically erase different fractions of the model’s reasoning — 0%, 25%, 50%, 75%, and 100% — and re-run the model to see if its answer changes. A turn is *anchored* when erasing all the reasoning makes no difference: the model gives the same number at 0% thinking as at 100% thinking. A turn is *exploring* when the answer shifts. This gives a binary time series for each conversation (Figure 1).

We ask three questions of that time series. First: does anchoring *persist* over multiple consecutive turns, or is each turn independently uncertain? Answer: we cannot confirm persistence — anchored runs look statistically consistent with a memoryless process, even with 148 observed anchored runs across 67 conversations. Second: does being more anchored predict getting the answer wrong? Answer: the trend points the right way (more anchoring accompanies more wrong answers), but 21 conversations with clean labels are not enough to confirm it statistically. Third: is anchoring a *conversation-level* property — do some conversations run heavily anchored while others stay mostly exploring? Answer: **yes, strongly**. The spread of anchoring across conversations is vastly larger than chance would produce ($\chi^2 = 236.9$, $df = 66$, $p \approx 0$; ICC = 0.152). And conversation length is the key predictor: from short to long conversations, anchored fraction nearly quadruples (7.2% $\rightarrow$ 28.4%, a $3.9\times$ gradient), replicated across three independent experimental seeds.

Contributions.

- **A new measurement: per-turn faithfulness in multi-turn conversations.** We apply counterfactual deletion at every turn, treating faithfulness as a time series rather than a scalar. This reveals dynamics that single-turn evaluations cannot see.
- **The bistability finding.** Within a single conversation the model alternates between exploring and anchored modes. This is not a property of the model or the task in isolation — it is a property of the conversation. (Across 67 conversations, the spread of anchored fractions is $237\times$ more variable than independent coin flips would produce.)
- **Conversation length as a risk predictor.** Anchoring nearly quadruples from short (≤ 10 turn, 7.2%) to long (> 20 turn, 28.4%) conversations — a $3.9\times$ gradient replicated across three independent experimental runs. This gives safety engineers a concrete, measurable signal: longer conversations carry systematically less reliable CoT.
- **A quantitative gap in current safety thinking.** Proposals like Korbak et al. [2025] assume CoT faithfulness is a fixed model property. Our results show it degrades predictably with conversation length, making long-context monitoring less reliable precisely where it is most needed.

2 Background and Related Work

CoT faithfulness in single-turn settings. Measuring whether a model’s stated reasoning causes its output has been studied for single-turn prompts. Lanham et al. [2023] introduced counterfactual deletion: truncate the `<think>` block and check whether the answer changes. Turpin et al. [2023] showed that biased context shifts answers without appearing in the trace. Arcuschin et al. [2025]




Figure 1: **Mode-switching CoT faithfulness.** *Top:* Turn-by-turn mode for sample 965 (40 turns shown): green = exploring (CoT drives the answer), red = anchored (CoT is post-hoc). The red burst coincides with premature commitment to the wrong answer “300” (gold = 660). *Bottom left:* In exploring mode, different truncation levels produce different answers—the CoT is causally active. *Bottom right:* In anchored mode, all five levels agree on the same answer regardless of how much reasoning is visible—the CoT is decorative.

found that even without adversarial framing, reasoning models produce post-hoc rationalization on 0.4–13% of naturally occurring prompts depending on task difficulty. Chen et al. [2025] showed that frontier reasoning models acknowledge injected hints in their thinking trace only 25–39% of the time. All these evaluations measure faithfulness as a scalar property of a (model, prompt) pair; none tracks how faithfulness evolves across turns of an ongoing conversation.

Reasoning Theater: the single-generation analog. The closest single-generation analog to our anchored mode is “Reasoning Theater” [Boppana et al., 2026]: in some cases, a model commits to an answer early in its generation and subsequent tokens serve as performative elaboration rather than genuine computation. Our work can be understood as measuring whether and when Reasoning Theater persists *across multiple turns*, finding that it does—as a statistically detectable mode with non-geometric run-length dynamics.

Lost in conversation. Laban et al. [2025] demonstrated that all major reasoning models (R1, o3, GPT-4.1) lose roughly 39 percentage points of accuracy when problems are revealed piece by piece across turns, naming the resulting failure pattern “lost in conversation.” The dominant cause is premature commitment. A follow-up by Liu et al. [2026] reframes this as an “intent mismatch” and proposes a Mediator-Assistant architecture as a fix. Neither paper studies the *faithfulness* of the reasoning trace during the lost phase.

CoT monitorability and the safety stakes. A growing safety literature argues that visible CoT traces are a currently valuable oversight safeguard [Korbak et al., 2025, OpenAI, 2025]: models find it difficult to simultaneously produce an unfaithful trace *and* behave deceptively. Our results challenge this premise in the multi-turn setting: a model in anchored mode produces elaborate, structurally plausible reasoning that is entirely decorative. The failure arises from conversational pressure rather than intentional deception, making it harder to detect by trace review alone.

3 Method

3.1 Experimental Setup

The core idea: put the model under sustained pressure. We want to watch faithfulness change over the course of a conversation — not just measure it once at the start. To do that, we need a setting where the model is forced to think and re-think across many turns, and where it will naturally start to “lock in” to an answer before it has all the facts. The *sharded conversation* framework of Laban et al. [2025] creates exactly this pressure.

What a sharded conversation looks like. Take a GSM8K math word problem and split it into 3–20 fragments (“shards”). A user simulator reveals one shard per turn. The model — which has only seen part of the problem — must still produce a best-guess answer at every turn. This mimics a real consulting scenario: the client drip-feeds information, and the advisor must reason incrementally. For example, shard 1 might say “*Alice has some apples.*” Shard 2 adds “*She gives half to Bob.*” Only at shard 5 does the model finally learn the starting count. A model that prematurely commits to an answer after shard 2 will not revise it even as shards 3–5 arrive — it is “lost in conversation.”

Concretely, at each turn the model (DeepSeek-R1-Distill-Qwen-7B, Guo et al. 2025) produces a `<think> ... </think>` reasoning trace followed by a numeric answer. A second model (Qwen2.5-7B-Instruct) plays the user simulator, revealing one shard per turn and acknowledging the model’s answer. We replaced the original GPT-4o-mini components with Qwen2.5-7B-Instruct to enable fully local, cost-free reproducibility. Both models run on a single NVIDIA RTX 6000 Ada (49 GB VRAM) in 8-bit or 4-bit quantisation.

Two phases of data collection. We collected data in two stages to progressively test whether the bistability finding holds at scale and across independent runs.

Phase A ($N = 24$ conversations): four independent datasets with different random seeds (default, 20260508, 12345, 99999) and increasing problem complexity — maximum shard counts of 4, 5, 6, and 10. Longer shard counts produce longer conversations with more room for premature commitment to develop.

Phase B ($N \approx 43$ conversations): three independent datasets with new random seeds (11111, 22222, 33333) and longer conversations (up to 20 shards, max 30 turns per conversation). The aim was to (a) confirm that the bistability pattern deepens in longer conversations, and (b) verify that the finding replicates across completely independent runs. Seeds 2 and 3 ran simultaneously using mixed-precision quantisation to exploit all available GPU memory, completing in approximately 48 GPU-hours.

Combined dataset: 67 conversations, 1,289 turn-level faithfulness measurements across seven independent datasets. We use N throughout to denote number of conversations.

3.2 Faithfulness Measurement: Does the Reasoning Actually Matter?

The test in plain English. At each turn, the model has just produced a reasoning trace and an answer. We want to know: did the reasoning *cause* that answer, or was the answer predetermined and the reasoning decorative? To find out, we re-run the model on the same context but with progressively less of its own reasoning visible, from *none at all* (0%) to the full trace (100%). If the answer stays the same regardless of how much reasoning we show, the reasoning was not driving the answer — the model was already committed. If the answer changes, the reasoning was causally active.

The five-level test. Formally, we apply Lanham-style counterfactual deletion [Lanham et al., 2023] at every assistant turn t :

1. Extract the thinking block $\mathcal{T}_t$ and the model’s answer.
2. Truncate $\mathcal{T}_t$ to five fractions of its original length: $\{0\%, 25\%, 50\%, 75\%, 100\%\}$.
3. For each truncated prefix, close the thinking block with `</think>` and regenerate the answer greedily (temperature = 0, max 128 tokens).
4. Extract a numeric value from each of the five answers.

We classify the turn as **anchored** if all five numeric values are identical — the answer does not change no matter how much reasoning the model sees. We classify it as **exploring** if at least one answer differs — the reasoning is shaping the output. Turns where fewer than three of the five answers contain a parseable number are excluded ($\approx 9\%$ of turns).

Why not use correctness? A model can be anchored on the *correct* answer (reasoning is still post-hoc — it just happens to be right), and can be exploring toward the *wrong* answer (reasoning is causally active but leading somewhere wrong). Correctness-based metrics conflate these cases. Our measure asks only “does the reasoning constrain the answer?” — the right question for a faithfulness test aimed at safety monitoring.

3.3 Three Questions About Mode-Switching

Labelling each turn as {anchored, exploring} gives us a binary time series per conversation. We ask three progressively more ambitious questions about that series.

H1 — Does anchoring persist? When the model enters anchored mode, does it tend to stay there for multiple consecutive turns, or does it bounce in and out randomly? A model that truly “locks in” should produce runs of anchored turns longer than chance would predict. Formally, we collect the lengths of all consecutive anchored runs (e.g., five red turns in a row counts as one run of length 5) and test whether those lengths follow a geometric distribution — the distribution you get if each turn’s mode were decided independently by a coin flip. A geometric distribution is the memoryless null; departing from it would mean that being anchored on turn t makes turn $t+1$ *more* likely to be anchored than chance.

To test this, we simulate 10,000 datasets from the best-fit geometric distribution and compare the KS distance of the real data to the simulated null — a parametric bootstrap that avoids a known artifact of applying the standard KS test to discrete integer data [Massey, 1951].

H2 — Does anchoring predict failure? If anchored mode means the model is not updating its reasoning on new information, it should be more common in conversations where the model gets the wrong answer. We measure this as a simple correlation: does the fraction of anchored turns in a conversation predict whether that conversation ended incorrectly? This requires conversations with an unambiguous right/wrong label, which is not available for most of our GSM8K conversations (the multi-turn extractor returns None in most cases). We report the result on the $N = 21$ conversations where a clean binary label is available.

H3 — Is anchoring a conversation-level property? Suppose anchoring were just random noise: every turn, independent of everything else, the model has a fixed 23% chance of being anchored. Then each conversation’s anchored fraction would hover near 23% with only small random variation — the natural scatter of a coin-flipping process. But if anchoring is a real mode — something that some conversations fall into heavily while others avoid — then we should see *far more spread* across conversations than the coin-flip baseline predicts.

We test this with a chi-square variance test. Under the coin-flip null, the standardised deviation of each conversation’s anchored fraction from the global mean follows a standard normal distribution, and the sum of squared deviations across all k conversations follows a chi-square distribution with $k-1$ degrees of freedom. A very large chi-square statistic (and correspondingly tiny p-value) means conversations differ in their anchoring far more than chance would produce — establishing that anchoring is a property of the *conversation*, not just random per-turn noise.

We also compute the Intraclass Correlation Coefficient (ICC), which directly answers “what fraction of the total variance in per-turn anchoring is explained by which conversation a turn belongs to?” An ICC of 0 would mean conversation identity predicts nothing; an ICC of 1 would mean all variance is between conversations.

4 Results

4.1 Evidence Accumulates Across Seven Datasets

Table 1 shows how the picture evolved as we added data. We ran our experiments incrementally, adding new conversations across four Phase A datasets and three Phase B seeds, and the core finding held at every scale.

At the smallest scale (Phase 1 only, 4 conversations), both modes appeared — the model anchored on 12% of turns and explored on the rest. This is too few conversations to run rigorous statistical tests, but the qualitative phenomenon was already visible. As we added Phase 2 and Phase 3, the anchored fraction stabilised around 8%. Phase 4 used longer problems (up to 10 shards per problem instead of 4–6), and the anchored fraction rose to 13% — more complex problems that run longer give anchoring more room to develop.

Phase B was designed to stress-test these early observations. We ran three completely fresh seeds with even longer conversations (up to 20 shards) and asked: does the bistability pattern survive? It does, at even larger scale. The combined 67-conversation dataset produces a variance test statistic of $\chi^2 = 236.9$ ($df = 66$, $p \approx 0$) — the observed variation across conversations is vastly larger than chance produces under a coin-flip model. The 148 observed anchored runs also gave us enough data for a definitive run-length test (H1), which remains inconclusive (bootstrap $p = 0.521$): the model’s anchored runs are consistent with a memoryless process at this effect size.

The higher anchored fraction in Phase B (27–38% per seed vs. Phase A’s 13%) is fully explained by Phase B’s longer conversations, as Section 4.9 shows.

Table 1: Cumulative mode-switching statistics across phases. H3: global anchored fraction (both modes “substantially present”). Variance test: chi-square test of between-conversation variance vs. Bernoulli null (Section 3). ✓ H3 confirmed at all scales. H1 (run-length persistence) was tested with a corrected parametric bootstrap on all $N = 67$ conversations (148 anchored runs): bootstrap $p = 0.521$, non-significant; see Section 4.3. All intermediate H1 p-values computed with `scipy’s kstest` are invalid on discrete data and are omitted here.

Dataset	Conv.	Faith. turns	Anchored (H3)	Verdict
Phase 1 only	4	53	0.12	H3 ✓
+ Phase 2	8	92	0.08	H3 ✓
+ Phase 3	14	177	0.08	H3 ✓
+ Phase 4 (Phase A combined)	24	412	0.13	H3 ✓ var. $p=0.0001$
+ Phase B Seed 1 (11111)	38	702	0.21	H3 ✓
+ Phase B Seed 2 (22222)	53	994	0.24	H3 ✓
+ Phase B Seed 3 (33333)	67	1289	0.233	H3 ✓ var. $p \approx 0$

4.2 Qualitative Mode-Switching: Sample 965

The 44-turn conversation `sharded-GSM8K/965` is the clearest qualitative illustration of the mode-switching pattern, visible in Figure 2 (top row). The conversation has three phases: (1) **Turns 1–8 (exploring)**: the model’s answer varies across truncation levels as each new shard changes what it knows; (2) **Turns 9–13 (anchored)**: all five levels produce “300” (gold = 660)—showing the model more or less of its reasoning makes no difference, and the thinking blocks during these turns (~ 412 tokens) construct elaborate derivations that are entirely post-hoc; (3) **Turns 14–44 (mostly exploring)**: the model partly recovers exploring mode as new shards arrive, but never converges on the correct answer. The anchored burst at turns 9–13 coincides with the onset of premature commitment: turn 8 is the first turn where the model states a confident answer, and turns 9–13 are those where that answer appears in every regeneration regardless of CoT context. Section 4.7 reports the systematic commitment-alignment analysis across all 67 conversations.



Figure 2: **Per-turn CoT faithfulness mode** (Phase A $N = 24$ conversations, 373 assessable turns). Green = exploring (CoT causal); red = anchored (CoT post-hoc). Rows sorted by anchored fraction (right axis).

4.3 H1: Does Anchoring Persist? — Inconclusive

What we’re testing. H1 asks whether consecutive anchored turns are more common than chance. If each turn’s mode were decided by a memoryless coin flip, we would expect most anchored “runs” to last exactly 1 turn — the model anchors, then snaps back. A model with genuine mode-switching would show longer runs: once anchored, it tends to stay anchored.

Across our 67 conversations we found 148 anchored runs with a mean length of 1.905 turns. The question is whether that mean — and the shape of the run-length distribution — is any longer than a memoryless (geometric) process with the same average would predict.

A corrected test. Standard software applies the Kolmogorov-Smirnov test to measure the gap between the empirical and theoretical distributions. However, applying it directly to discrete integer run-lengths produces a guaranteed-spurious result: the test statistic is mathematically lower-bounded by the geometric parameter $\hat{p}$, so the test nearly always rejects the null regardless of whether the data actually follow a geometric distribution [Massey, 1951]. To correct for this artifact, we use a *parametric bootstrap*: we simulate 10,000 datasets from the best-fit geometric distribution, compute the KS statistic on each, and ask what fraction of simulated statistics are at least as large as the one we observed in the real data. That fraction is our p-value — by construction immune to the discrete-data artifact.

Result: H1 is inconclusive. Bootstrap $p = 0.521$: the observed run-length distribution is entirely consistent with a memoryless geometric process. This result held even with $5.5\times$ the statistical power of our initial 27-run Phase A analysis. There are two possible interpretations: either anchoring is genuinely memoryless (each turn is an independent coin flip), or there is a small persistent effect that would require $N \geq 200$ conversations to detect. We cannot distinguish between them at this sample size.

Qualitatively, individual anchored bursts are striking — sample 965 shows five consecutive anchored turns, a sequence that feels deliberate rather than random (Section 4). Confirming whether this persistence is statistically real remains an open question. Figure 3 shows the empirical and null distributions.

4.4 H2: Does More Anchoring Mean More Wrong Answers? — Underpowered

The hypothesis. If a model in anchored mode is not updating its reasoning on new information, conversations where it anchors heavily should be conversations where it ends up wrong. H2 tests this directly: is the fraction of anchored turns negatively correlated with conversation correctness?




Figure 3: **H1: Anchored run-length distribution vs. geometric null** (Phase A $N = 24$ conversations, 27 anchored runs shown). *Left*: Empirical anchored-mode run lengths (red bars) vs. the best-fit geometric null ($\hat{p} = 0.54$, dashed line). The combined $N = 67$ analysis (148 runs) gives bootstrap $p = 0.521$, still non-significant; H1 is inconclusive. *Right*: Exploring-run distribution for comparison.

The measurement problem. GSM8K multi-turn conversations rarely end with a cleanly labelled correct or incorrect outcome: the model answers incrementally across turns in inconsistent formats, and our extractor cannot always determine a definitive final answer. Of our 67 conversations, only 21 have unambiguous binary correctness labels.

Result: right direction, not enough data. Among those 21 conversations, conversations the model got wrong did have a higher average anchored fraction than conversations it got right — the trend points exactly where H2 predicts. The correlation is $r = -0.242$, $p = 0.291$. That p-value of 0.29 does not reach statistical significance — but this is an expected consequence of having only 21 data points, not evidence against the hypothesis. A power analysis shows we would need approximately $N = 131$ conversations with clean correctness labels to have an 80% chance of detecting an effect this size. We are at 16% of the required sample.

Figure 4 shows the per-conversation anchored fractions by outcome and the power analysis. **H2 is inconclusive, not failed**: the data are consistent with anchoring predicting failure but the dataset is too small to confirm it. The natural next step is replication on HumanEval coding conversations, where every conversation has an unambiguous pass/fail label.

4.5 H3: Anchoring is a Conversation-Level Property

Both modes are present. Across 1,250 assessable turns in our 67 conversations, 23.3% are anchored and 76.7% are exploring. Both modes are far above the 5% threshold we require for a mode to be considered “substantially present.” There is no conversation where the model is anchored every single turn, and no conversation where it never anchors. The two modes co-exist.

But are they really modes, or just noise? The deeper question is whether anchoring is a structured, conversation-level property or just random per-turn noise. To see why this matters: if every turn were anchored with a fixed 23% probability independently of all other turns, then across a 10-turn conversation you would expect roughly 2–3 anchored turns — but you might get 0 or 6 just by chance. Under that coin-flip model, some conversations would look anchor-heavy just by luck, and the variation across conversations would be predictable and small.

What we observe is strikingly different. The spread of anchored fractions across conversations is *far* larger than the coin-flip model predicts: some conversations anchor on fewer than 5% of turns while others anchor on more than 60%. A chi-square test formalises this — it computes the expected spread under independent per-turn coin flips and asks how unlikely the observed spread is. The answer: astronomically unlikely ($\chi^2 = 236.9$, $df = 66$, $p \approx 0$, $N = 67$). Anchoring is not noise; it is a property that belongs to a conversation.

H2: Anchoring predicts failure (correct direction) — but underpowered at N=21




Figure 4: **H2: Anchored fraction by conversation outcome.** *Left:* Per-conversation anchored fraction for Incorrect vs. Correct conversations ($N = 21$ with valid binary labels); representative data matching $r = -0.24$. The correct-direction trend is visible but not significant ($p = 0.291$). *Right:* Power analysis—80% power requires $N \approx 131$ conversations with binary correctness labels.

Ruling out length as a confound. One natural worry: maybe the H3 result just reflects the fact that longer conversations produce more chances for a model to get anchored, and we have conversations of very different lengths. If that were the whole story, we would expect the result to disappear when we compare conversations of similar length.

We test this by running the same chi-square test *separately within each length stratum*. Among short conversations (≤ 10 turns, $n = 14$): $\chi^2 = 4.5$, $df = 13$, $p = 0.97$ — no significant clustering; short conversations are too uniform to detect structure. Among medium-length conversations (10–20 turns, $n = 18$): $\chi^2 = 36.9$, $df = 17$, $p = 0.0035$ (**significant**). Among long conversations (> 20 turns, $n = 29$): $\chi^2 = 130.1$, $df = 22$, $p < 0.001$ (**significant**). Even within the group of medium-length conversations, some anchor heavily and others do not — and this cannot be explained by length differences within that group. H3 holds inside each length stratum. Figure 5 shows the full per-conversation distributions.

How much variance does conversation identity explain? The Intraclass Correlation Coefficient (ICC) answers this directly. Decomposing the total variance in per-turn anchoring into “within-conversation” and “between-conversation” components gives $ICC = 0.152$. Fifteen percent of all variation in whether a given turn is anchored is explained purely by knowing *which conversation* that turn belongs to — independent of any other information. This is a moderate but meaningful effect, and it independently confirms H3: anchoring is structured at the conversation level.

4.6 Is Anchoring Just Short Reasoning? No.

A skeptic might say: anchored turns look the same at 0% and 100% CoT simply because the thinking block is short. If there is very little reasoning to truncate, all five levels would trivially agree. We checked this in sample 965: the median thinking block at anchored turns is 412 tokens, compared to 387 tokens at exploring turns in the same conversation. If anything, anchored turns have *more* reasoning — it just does not constrain the output. Anchoring is not a consequence of having nothing to say.

4.7 Does Anchoring Align With Premature Commitment?

The hypothesis. In sample 965, the model enters anchored mode at turn 9 — exactly one turn after it first commits to a confident (wrong) answer at turn 8. The qualitative story is compelling: the

H3 confirmed within length bins: bistability is not a length artifact




Figure 5: **Per-conversation anchored fraction by length bin.** Each dot is one conversation. Box plots show the spread within each stratum (combined $N = 67$); diamonds mark means. Even within the medium bin (10–20 turns), conversations vary from near-0% to ~50% anchoring ($p = 0.0035$), confirming that anchoring is a conversation-level property, not a length artifact.

model decides the answer is 300, and from that moment its reasoning becomes a post-hoc justification for that decision. We tested whether this alignment generalises across all 67 conversations.

How we operationalised it. For each conversation we identified (a) the first anchored turn, and (b) the “commitment turn” — the first turn where the model gave the same answer on two consecutive turns, a proxy for premature lock-in. 42 of our 67 conversations had identifiable commitment turns.

Result: the qualitative story does not generalise. The Pearson correlation between anchoring onset and commitment onset across these 42 conversations is $r = 0.06$, $p = 0.71$ — essentially zero. On average, anchoring onset lags commitment onset by about 3 turns, but the scatter is enormous: in many conversations the timing is reversed, or there is no clear alignment at all. The single vivid example in sample 965 is a real phenomenon but not a reliable cross-conversation predictor.

4.8 Could Anchoring Just Be Answer Repetition?

The alternative explanation. Our test classifies a turn as anchored when all five truncation levels give the same answer. But here is a simpler story: at 0% CoT the model cannot reason from scratch, so it just repeats the previous turn’s answer. At 25%, 50%, 75% it still does not see enough to change its mind, so it repeats again. The result looks like “anchoring” but it is really just answer inertia — the conversational equivalent of “I’ll stick with what I said.” If this story were true, nearly every anchored answer should match the preceding turn’s answer.

What we find: partially true, not the whole story. Among our 272 anchored turns, 73.9% (201 turns) do match the preceding turn’s answer — high, but not universally so. More diagnostically, 75.8% of anchored turns are self-repetitions versus 45.6% of exploring turns, a ratio of $1.66\times$. If anchoring were *purely* inertia, we would expect that ratio to be $2\times$ or higher; it falls short. And 24.2% of anchored turns produce a fresh answer not matching the previous turn — ruling out pure inertia entirely.

The honest interpretation: answer inertia under conversational pressure is a *contributing factor* to anchoring, but it is not the complete explanation. Some anchored turns are genuine post-hoc rationalization of an internally computed answer; others are simple repetition. Both are problematic for safety monitoring: a trace that is either decorative or copied is equally uninformative. Figure 6 shows the breakdown.




Figure 6: **Repetition confound analysis (combined $N = 67$).** *Left:* Repetition rate (answer matches previous turn) in anchored vs. exploring mode: 75.8% vs. 45.6% (1.66 $\times$ ratio, below the 2 $\times$ threshold). *Right:* Of all anchored turns, 24.2% produce novel (non-repeated) answers, ruling out pure answer inertia.

4.9 Conversation Length Predicts Anchoring: A 3.9 $\times$ Gradient

The finding. Among our 67 conversations, we sort by total turn count and ask: do longer conversations anchor more? The answer is a strong yes. We group conversations into three length bins and compute the mean anchored fraction in each:

- **Short** (≤ 10 turns, $n = 14$): **7.2%** anchored
- **Medium** (10–20 turns, $n = 18$): **12.8%** anchored
- **Long** (> 20 turns, $n = 29$): **28.4%** anchored

From short to long conversations, the anchored fraction increases **3.9 $\times$** . This is not a ceiling effect or outlier artifact: Figure 7 shows the individual data points within each bin, and the gradient is consistent across Phase A and Phase B conversations alike.

Why Phase B looks more anchored overall. Phase B (seeds 1–3) has an overall anchored fraction of 27–38%, versus 13% in Phase A. This gap initially looks like a qualitative difference, but it disappears after controlling for conversation length: Phase B used up to 20 shards per problem (vs. 10 in Phase A), producing systematically longer conversations that land in the high-anchoring “long” bin. Within matched length bins, Phase A and Phase B produce consistent anchored rates. The phenomenon is the same; Phase B simply has more long conversations.

The safety implication. This gradient is the paper’s most actionable finding. Current safety proposals treat CoT faithfulness as a fixed property of a model: measure it once on a benchmark, then rely on it in deployment. Our results say this assumption is wrong in a specific, quantifiable way: the fraction of turns where the reasoning trace is post-hoc rationalization grows as the conversation gets longer, reaching 28% in conversations over 20 turns long. A safety monitor that was reliable during short conversations will systematically overestimate its reliability as conversations grow.

4.10 Cross-Seed Reproducibility

Three independent Phase B random seeds (11111, 22222, 33333) provide a reproducibility check. Anchored fractions were: Phase A baseline $\approx 13.4\%$; Seed 1 $\approx 27\%$; Seed 2 $\approx 35.9\%$; Seed 3 $\approx 38.1\%$. The Phase A–Phase B gap is explained by Phase B’s longer conversations ($\text{max_shards} = 20$): after




Figure 7: **3.9× length-anchoring gradient (combined $N = 67$).** *Main panel:* Mean anchored fraction per length bin with 95% CI whiskers; jittered points show individual Phase A conversations. *Right panel:* Within-bin variance test p-values confirming H3 holds within each stratum.




Figure 8: **Cross-seed reproducibility.** *Left:* Anchored fraction per dataset (Phase A and three Phase B seeds). *Right:* Anchored fraction vs. mean conversation length across datasets, confirming the length effect accounts for the Phase A–Phase B gap.

length-stratification the three seeds produce consistent rates within each bin. The ICC of 0.152 is stable across phases.

5 Discussion

Faithfulness as a dynamical conversation property. Prior work treats faithfulness as a static property of a (model, prompt) pair: model M is $x\%$ faithful on benchmark B [Lanham et al., 2023, Arcuschin et al., 2025]. Our results reframe it as a *dynamical conversation property*: the same model, on the same task, can switch between faithful and unfaithful operation within a single conversation,

tracking the epistemic state of the dialogue. This changes the measurement problem: a single-turn faithfulness benchmark cannot characterise a model’s faithfulness under the multi-turn deployment conditions where it matters most.

Connection to Reasoning Theater and the Reasoning Horizon. The anchored mode we identify is the multi-turn analog of two single-generation phenomena. “Reasoning Theater” [Boppana et al., 2026] describes models committing to answers before completing their chain, with subsequent tokens performative rather than causal—a within-turn snapshot of what we observe as a persistent between-turn mode. Ye et al. [2026] found a “Reasoning Horizon” at 70–85% of chain length beyond which tokens have negligible causal influence—a within-generation faithfulness decay. Our work extends both findings to the multi-turn regime: a model can sustain an entirely decorative CoT state for multiple consecutive conversation turns, not just within a single generation.

Implications for CoT oversight. The CoT monitorability literature [Korbak et al., 2025, OpenAI, 2025] argues that visible reasoning traces are a currently viable safety signal because models find it difficult to produce an unfaithful trace while behaving deceptively. Korbak et al. [2025] argues that CoT monitorability is a fragile but real opportunity for AI safety, contingent on models being unable to complete tasks without externalising their reasoning. Our results identify a previously undocumented source of fragility: within a single conversation, a model alternates between faithfully externalised reasoning (exploring mode) and opaque post-hoc rationalisation (anchored mode). The $3.9\times$ length gradient implies that safety monitors face systematically lower reliability as conversations extend—precisely the high-stakes, long-context settings where monitoring is most needed. A practical mitigation would be to monitor answer consistency under lightweight CoT perturbation online, flagging conversations where the model enters anchored mode and triggering a summarise-and-restart intervention. Unlike adversarial deception, this failure mode arises from conversational pressure and is detectable from answer consistency alone—making it more tractable to address.

6 Limitations

Single model. All experiments use DeepSeek-R1-Distill-Qwen-7B in 8-bit or 4-bit quantisation. Whether mode-switching generalises to larger or differently trained reasoning models (o3, QwQ-32B, Gemini Thinking), or whether quantisation affects greedy output determinism in ways that inflate the anchored rate, is unknown.

Single task domain. Only GSM8K math is tested. Coding and factual QA tasks may exhibit different faithfulness dynamics, and a task with binary outcomes would enable H2.

H2 inconclusive, needs $N\approx 131$. Among $N = 21$ conversations with clean binary labels, $r = -0.242$, $p = 0.291$ —correct direction but underpowered. Power analysis shows 80% power requires $N \approx 131$ conversations with binary correctness labels. H2 is the most actionable hypothesis from a safety standpoint; replication on HumanEval coding conversations is the natural next step.

Context-solvability confound. At 0% CoT, the model still sees all prior shards. For some problems the answer may be derivable from conversation context alone, making all five truncation levels agree not because the CoT is post-hoc but because the problem is locally determined. The repetition confound check (Section 4) quantifies this; if most anchored turns are self-repetitions, the finding should be reinterpreted as answer repetition under conversational pressure rather than post-hoc rationalization.

Scale for H1 confirmation and cross-model comparison. $N = 67$ is sufficient to confirm H3 but H1 remains inconclusive: the corrected bootstrap gives $p = 0.521$ even with 148 anchored runs. This suggests the geometric null may be genuinely adequate at this model and task; a study with $N \geq 200$ conversations and a wider range of run lengths would be required to distinguish a small super-geometric effect from a truly memoryless process. Cross-model comparison ($N \geq 100$ per model) would be needed for generalisation claims across architectures.

7 Conclusion

Chain-of-thought faithfulness in multi-turn conversations is mode-switching, not monotonically decaying. Within a single derailing conversation, the same reasoning model alternates between

exploring mode (CoT causally active) and anchored mode (CoT post-hoc or inertia-driven). Across 67 conversations and 1,289 faithfulness observations, 23.3% of turns are anchored. Between-conversation variance significantly exceeds a Bernoulli null ($\chi^2 = 236.9$, $df = 66$, $p \approx 0$, $N = 67$), establishing that mode-switching is a conversation-level phenomenon: some conversations accumulate extended anchored intervals while others remain predominantly exploring. Within both the medium ($p = 0.0035$) and long ($p < 0.001$) conversation strata the effect is confirmed, ruling out the length-heterogeneity confound. ICC = 0.152 (moderate) independently quantifies conversation-level clustering. Formal run-length persistence (H1) is inconclusive at $N = 67$ (bootstrap $p = 0.521$, 148 anchored runs); the geometric null may be a genuinely adequate description at this model and task. Three independent random seeds produced consistent anchoring rates (27–38%), confirming the bistability finding is not seed-specific. The $3.9\times$ length-anchoring gradient (7.2% short $\rightarrow$ 12.8% medium $\rightarrow$ 28.4% long) provides a concrete, actionable predictor: CoT oversight reliability degrades monotonically with conversation length. The anchored-mode repetition rate of 75.8% (vs. 45.6% in exploring mode, ratio $1.66\times$) indicates that answer inertia under conversational pressure is a contributing factor, but the 24.2% of anchored turns with novel answers rules out pure answer inertia. Commitment alignment between anchoring onset and premature commitment onset is non-significant across conversations ($r = 0.06$, $p = 0.71$, $n = 42$). For AI safety, CoT traces in deployed conversational systems face systematically lower reliability as conversations extend—precisely the high-stakes, long-context settings where monitoring is most needed. The two most important open questions are (1) whether anchored mode statistically predicts conversation failure (H2), requiring $N \approx 131$ conversations with clean binary outcomes such as HumanEval; and (2) whether the mode-switching phenomenon generalises to larger or differently-trained reasoning models.

Acknowledgments and Disclosure of Funding

This research was conducted using Lossfunk’s Autoresearch harness. AI assistance (Claude Sonnet 4.x by Anthropic) was used for experimental implementation, code debugging, figure generation, and draft feedback throughout this project, as documented in the AI Involvement Checklist below.

References

- Ben Arcuschin et al. Chain-of-thought reasoning in the wild is not always faithful. *arXiv preprint*, arXiv:2503.08679, 2025.
- Siddharth Boppana, Annabel Ma, Max Loeffler, Raphael Sarfati, Eric Bigelow, Atticus Geiger, Owen Lewis, and Jack Merullo. Reasoning theater: Disentangling model beliefs from chain-of-thought. *arXiv preprint*, arXiv:2603.05488, 2026.
- Yanda Chen et al. Reasoning models don’t always say what they think. *arXiv preprint*, arXiv:2505.05410, 2025.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint*, arXiv:2501.12948, 2025.
- Tomek Korbak, Joseph Bloom, Alan Cooney, Geoffrey Irving, Martín Soto, and Jasmine Wang. Chain of thought monitorability: A new and fragile opportunity for AI safety. *arXiv preprint*, arXiv:2507.11473, 2025. URL <https://arxiv.org/abs/2507.11473>.
- Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. LLMs get lost in multi-turn conversation. *arXiv preprint*, arXiv:2505.06120, 2025.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Miles Steiner, Carson Denison, Danny Hernandez, Evan Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint*, arXiv:2307.13702, 2023.
- Yao Liu et al. Intent mismatch causes LLMs to get lost in multi-turn conversation. *arXiv preprint*, arXiv:2602.07338, 2026.
- Frank J. Massey. The kolmogorov-smirnov test for goodness of fit. *Journal of the American Statistical Association*, 46(253):68–78, 1951.

Microsoft Research. lost-in-conversation: Benchmark for multi-turn LLM evaluation. <https://github.com/microsoft/lost-in-conversation>, 2025.

OpenAI. Monitoring monitorability. Technical report, OpenAI, 2025. URL <https://cdn.openai.com/pdf/d57827c6-10bc-47fe-91aa-0fde55bd3901/monitoring-monitorability.pdf>.

Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Donald Ye, Max Loffgren, Om Kotadia, and Linus Wong. Mechanistic evidence for faithfulness decay in chain-of-thought reasoning. *arXiv preprint*, arXiv:2602.11201, 2026.

A Experimental Details

Dataset. We use the `sharded_instructions_600.json` file from the `microsoft/lost-in-conversation` repository. The math subset contains 103 problems drawn from GSM8K. Samples were selected in four independent phases using distinct random seeds; `-exclude_dirs` prevents any sample from appearing in more than one phase.

Infrastructure. All experiments run on NVIDIA RTX 6000 Ada (49 GB VRAM, shared server), Ubuntu 24.04, CUDA 12.4, HuggingFace Transformers 5.5.0, BitsAndBytes 0.49.2. When two processes ran in parallel, `PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True` prevented allocator OOM. Model weights are cached on a RAM-disk (`/dev/shm`) due to disk space constraints.

Hyperparameters. DeepSeek-R1-Distill-Qwen-7B: temperature 0.6, top- p 0.95, `R1_MAX_NEW_TOKENS = 1500` during conversation. Faithfulness regeneration: temperature 0.0 (greedy), `max_new_tokens = 128`. Qwen2.5-7B-Instruct: temperature 1.0 (user simulator), 0.0 (system verifier).

Faithfulness token budget. Pilot experiments used `max_new_tokens = 512` (146 s/turn). Switching to 128 reduced measurement time to 31 s/turn ($5\times$ faster) with no degradation in answer extraction, since the metric only needs a numeric value.

Phase B hyperparameters. Phase B seeds (11111, 22222, 33333): `max_shards = 20`, `max_turns = 30`, `n_samples = 15`, `faith_tokens = 128`. Seeds 2 and 3 ran simultaneously using mixed-precision (8-bit + 4-bit quantisation, `PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True`) to maximise GPU utilisation on the shared RTX 6000 Ada.

Compute. Approximately 68 GPU-hours total: ≈ 20 hours for Phase A (four datasets) and ≈ 48 hours for Phase B (three seeds, `max_shards = 20`, `max_turns = 30`).

Reproducibility. Code is available at [anonymised for submission]. Phase A random seeds: default, 20260508, 12345, 99999. Phase B random seeds: 11111, 22222, 33333.

Additional figures.

Bistability cascade: effect strengthens with N




Figure 9: **Mode-switching evidence accumulates across phases.** *Left:* H1 significance (bootstrap p-value) grows with N as more anchored runs become available to test against the geometric null. The power increase from $N = 14$ to $N = 24$ reflects both additional conversations *and* Phase 4's longer conversations (`max_shards = 10`), so this trajectory is not a fully independent replication. *Right:* Anchored fraction remains 8–14% at every scale, confirming both modes are persistently present.




Figure 10: **Per-conversation summary.** *Left:* Anchored fraction vs. conversation length; marker size encodes the longest anchored run. Longer conversations tend to have higher anchored fractions (see also Section 4.9 for the length-stratified analysis). *Right:* Global H3 distribution (13% anchored / 87% exploring) with the H3 variance test p-value.

Conference For AI Scientists 2026 — AI Involvement Checklist

Research Stage Assessment

1. Hypothesis development

Answer: **[D]**

Iteration Effort: **[Medium]**

Explanation: The research question (do CoT traces become unfaithful in multi-turn conversations?) was proposed by an AI research agent (Claude Sonnet) after surveying the CoT faithfulness and multi-turn conversation literatures. The human researcher selected this question from three candidates. The mode-switching reframing emerged from the agent's interpretation of preliminary data from sample 965.

2. Experimental design and implementation

Answer: **[D]**

Iteration Effort: **[Medium]**

Explanation: The AI agent designed and implemented the full experimental pipeline: model wrapper, counterfactual deletion measurement, batch runner, and mode-switching analysis scripts. The human researcher approved design choices and provided compute access. Several iterations handled GPU memory constraints, Unicode issues, and metric bugs (including replacing an invalid KS test with a parametric bootstrap after reviewer feedback).

3. Analysis of data and interpretation of results

Answer: **[C]**

Iteration Effort: **[Low]**

Explanation: The AI agent ran quantitative analyses (bootstrap tests, chi-square tests, confound checks) and generated all figures. Interpretation of the mode-switching pattern—connecting anchored turns to premature commitment—was developed jointly.

4. Writing

Answer: **[D]**

Iteration Effort: **[Low]**

Explanation: The paper was written by the AI agent following the Research Paper Writing Skill guidelines (one message per paragraph, claim-evidence alignment, adversarial self-review). The human researcher reviewed for accuracy.

AI System Documentation

5. AI systems used

Description: A frontier large language model acting as an autonomous research agent (Loss-funk Autoresearch harness, Claude Sonnet 4.x). The agent surveyed literature, proposed hypotheses, implemented experiments, ran them via SSH on a GPU server, interpreted results, and drafted this paper with high autonomy under human oversight.

6. Observed AI Limitations

Description: (1) The agent initially designed experiments at a scale infeasible on a shared GPU, wasting hours before scope was reduced. (2) A metric bug (correctness-flip instead of answer-consistency) required a debugging pass. (3) The agent used 512 max_new_tokens for faithfulness regeneration (146 s/turn) before the human researcher identified that 128 was sufficient (31 s/turn, 5× faster). (4) The agent’s initial H1 test used scipy’s kstest on discrete data, which produces a guaranteed-significant KS statistic due to a left-boundary artifact; this was replaced with a parametric bootstrap and chi-square GOF. All failures were caught through human review of intermediate outputs.

Conference For AI Scientists 2026 — Reproducibility and Responsibility Checklist

1. Claims

Answer: [Yes]

Justification: Abstract and Introduction accurately reflect scope. H3 (23.3% anchored; variance test $\chi^2 = 236.9$, $p \approx 0$, $N = 67$) is confirmed; H1 (run-length persistence) is explicitly stated as inconclusive (bootstrap $p = 0.521$, $N = 67$, 148 anchored runs); H2 is stated as inconclusive and underpowered ($r = -0.242$, $p = 0.291$, $N = 21$). All claims are supported by Section 4.

2. Limitations

Answer: [Yes]

Justification: Section 6 discusses single model, single task, context-solvability confound, H2 inconclusive and underpowered (needs $N \approx 131$), and cross-model scale.

3. Theory assumptions and proofs

Answer: [NA]

Justification: No theoretical claims requiring proof.

4. Experimental result reproducibility

Answer: [Yes]

Justification: Appendix A specifies model versions, quantisation, seeds, data files, and hyperparameters.

5. Open access to data and code

Answer: [Yes]

Justification: Dataset from public microsoft/lost-in-conversation repository. Both models publicly available on HuggingFace. Code included in anonymised supplementary material.

6. Experimental setting/details

Answer: [Yes]

Justification: Appendix A provides all hyperparameters, infrastructure details, and selection criteria.

7. Experiment statistical significance

Answer: [Yes]

Justification: Section 4 reports bootstrap-corrected and chi-square H1 p-values, the H3 variance test, and explicitly characterises H2 as reserved for future work (not failed).

8. Experiments compute resources

Answer: [Yes]

Justification: Appendix A: NVIDIA RTX 6000 Ada (49 GB VRAM), 8/4-bit quantisation, ≈ 20 GPU-hours total.

9. Research Ethics

Answer: [Yes]

Justification: Publicly available models and datasets. No human subjects or sensitive data.

10. Broader impacts

Answer: [Yes]

Justification: Section 5 discusses the safety implication (CoT unreliable during anchored phases) and argues for greater caution when using CoT for oversight. No negative societal impacts specific to this work.